

Framework Comparison for Legal RAG Systems: LangChain vs LlamalIndex

Systematic empirical comparison of LangChain and LlamalIndex on the CUAD-RAG legal benchmark

Presented by: Triyasha Ghosh Dastidar
(tg2936)

Motivation & Problem Definition

Why RAG for legal NLP

- Domain-specific knowledge and factual grounding for contracts
- Reduces hallucinations; improves answer fidelity
- Handles complex terminology and long contracts

Why framework orchestration matters

- Retrieval effectiveness (recall@k) drives context quality
- Tool retrieval/invocation and agent reasoning affect reliability
- Latency/overhead trade-offs impact deployability

Current gaps in evaluations

- Lack of systematic framework-level comparisons
- Metrics often overlook tool accuracy and orchestration overhead
- No unified, reproducible pipeline integrating retrieval, generation, tools

Problem statement

How do LangChain and LlamaIndex compare for legal RAG under identical datasets, retrievers, LLMs, embeddings, and evaluation protocols?

System Overview (3A): Architectures & Models

Indexing

CUAD-RAG chunks + embeddings

Retrieval

Vector search (recall@k)

Generation

LLM answers

Tool Orchestration

Agent selects tools (if enabled)

Architectures evaluated

- Simple RAG: retrieval + single-step generation
- Tool-use Agent RAG: agent controller invokes specialized tools

Goal

Isolate framework orchestration behavior under identical models, embeddings, prompts, and evaluation protocols.

Models & embeddings

GPT-3.5-turbo

Used for Simple RAG (temperature = 0)

GPT-4o-mini

Used for tool-use agents (temperature = 0)

text-embedding-3-large

Embeddings for chunk vectors (shared across frameworks)

Note: Prompts, retrievers, and model settings held constant to ensure comparability.

Dataset & Indexing (3B): CUAD → CUAD-RAG

CUAD → CUAD-RAG transformation

 CUAD
41 clause categories



 Queries
Natural-language questions



 Gold spans
Mapped to contract text

- One query per contract; multiple gold spans retained
- Ambiguous/multipart contracts filtered to enforce one-to-one mapping

Chunking strategy

500 tokens

50 overlap

- LangChain: **RecursiveCharacterTextSplitter**
- LlamaIndex: **SentenceSplitter**

Vector stores & similarity

- LangChain → FAISS (cosine similarity)
- LlamaIndex → native vector store (cosine similarity)
- Shared embeddings: **text-embedding-3-large**

Test set & document characteristics

- 1,000 queries (standardized test set)
- Contracts range **5,000–50,000 tokens**
- Many queries require multi-location answer synthesis

Orchestration & Evaluation (3C)

Frameworks compared

LangChain

Modular chain-based design; LCEL

Llamaindex

Sophisticated retrieval & agent orchestration

Agent tools

📅 extract_dates Temporal extraction

🔍 rag_qa Retrieval-based QA

Retrieval depth sweep

k = 10

k = 20

k = 100

Evaluation framework

Answer quality

- F1, BLEU
- ROUGE-1/2/L
- BARTScore (semantic)

Retrieval

- Recall@k

Efficiency

- Latency (avg per query)
- Tool efficiency (expected vs invoked)

Tool usage

- Call counts & patterns

Experiment tracking

Weights & Biases (per-query and aggregate logs, execution traces)

Framework Comparison: LangChain vs LlamaIndex

Context: legal RAG on CUAD-derived benchmark ($k \in \{10, 20, 100\}$; tools: rag_qa, extract_dates)

Aspect	LangChain	LlamaIndex
Architecture	Chain/Agent/Tool primitives; prompt-first orchestration Flexible composition, rich integrations	Index → Retriever → QueryEngine graph Document-centric RAG with node abstractions
Retrieval	Pluggable retrievers; simple vector search by default Easy swapping and rerank via integrations	Node parsers, query transforms, hybrid/fusion Stronger recall@k in our setup
Tool orchestration	Mature agent patterns (ReAct, tool/function calling) Wide tool ecosystem; lower coordination overhead	Tools as QueryEngines/ToolSpecs; index routing Strong for RAG-centric tools and routing tasks
Performance (this project)	Lower latency on average Similar F1 despite lower recall at high k	Higher recall@k; better semantic scores (e.g., BARTScore/ROUGE) Slightly higher latency due to query transforms

Prefer LangChain when:

- Latency is critical and tool use is rich/heterogeneous
- You need broad integrations and rapid prototyping

Prefer LlamaIndex when:

- Retrieval quality and index routing are primary
- You want stronger semantic faithfulness in RAG

4A. Quantitative Results

Comparison on 1,000 CUAD-RAG queries ($k \in \{10, 20, 100\}$)

Dataset: CUAD-RAG Benchmark

System Performance Metrics

Metric breakdown by Retrieval Depth (k)

Framework	K	Recall@K	Latency (s)	Tool Eff.	F1 Score
LangChain	10	0.309	4.89	0.841	0.317
	20	0.342	4.84	0.873	0.335
	100	0.342	4.86	0.876	0.335
LlamaIndex	10	0.516	8.49	0.935	0.319
	20	0.656	9.52	0.950	0.341
	100	0.795	30.88	0.957	0.333

Semantic Quality (at k=20)

Metric	LangChain	LlamaIndex
BLEU	0.301	0.308
ROUGE-L	0.405	0.413
BARTScore	-1.189	-1.156

Key Performance Deltas

- ↗ **Retrieval Superiority:** LlamaIndex achieves **2.3x higher recall** at k=100 compared to LangChain.
- ⌚ **Latency Stability:** LangChain maintains stable **~4.8s latency** regardless of retrieval depth, whereas LlamaIndex spikes to 30s.
- ❖ **Semantic Alignment:** Despite similar F1 scores, semantic metrics (BARTScore, ROUGE) consistently favor LlamaIndex.

4B. Key Insights

Trade-off Analysis & Architectural Implications

 Analysis based on 1,000 test cases



Retrieval vs. Latency Trade-off

A fundamental trade-off exists between completeness and speed. LlamaIndex prioritizes **depth** (high recall) at the cost of exponential latency growth (up to 30s). LangChain optimizes for **stability** (~4.8s) but sacrifices retrieval coverage.

↳ Data: Recall +48% vs Latency +535% (LlamaIndex k=100)



The "F1 Illusion" & Semantics

While token-level F1 scores are comparable (~0.33), they mask quality differences. Better retrieval directly correlates with higher **semantic alignment** (BARTScore/ROUGE), meaning LlamaIndex answers are factually closer to gold standards despite similar keywords.

= Insight: Retrieval quality ≠ Generation tokens, but = Meaning



Retrieval Scaling Dynamics

Frameworks respond differently to increased context. LangChain's recall **plateaus** after k=20, suggesting architectural limits in candidate processing. LlamaIndex continues to gain recall up to k=100, effectively utilizing the expanded context window.

↗ Trend: LangChain saturates early; LlamaIndex scales



Orchestration Efficiency

LlamaIndex demonstrates higher **tool efficiency** (0.95 vs 0.87), indicating more precise agentic reasoning. It invokes tools slightly more often but with significantly better alignment to user intent, reducing wasted computational cycles on irrelevant actions.


◎ Stat: +8% Efficiency gap favors LlamaIndex

5. Demo & Visualization

Performance dynamics and agent behavior analysis


Live Metrics from WandB

Recall vs. Latency Trade-off




Notice: LlamaIndex (Indigo) pushes right (latency) for height (recall), while LangChain (Teal) clusters in the low-latency zone.

Agent Tool Orchestration (Avg Calls/Query)




Notice: Both agents favor rag_qa (~90%). LlamaIndex invokes tools slightly more frequently as context deepens.

Retrieval Recall@k Scaling



Notice: LangChain plateaus at k=20, whereas LlamaIndex continues to benefit from deeper retrieval (up to k=100).

Semantic Answer Quality (at k=20)



Notice: While F1 is similar, LlamaIndex leads in ROUGE-L and BLEU, indicating better structural answer alignment.

6. Lessons Learned

Implications for Legal RAG Design & Evaluation

💡 Key Takeaways from CUAD-RAG Study



Framework Choice Matters

Orchestration is not merely an implementation detail. It fundamentally dictates **retrieval depth**, **latency profiles**, and **tool behavior** beyond simple model accuracy.

✓ **Architectural Trade-offs**



Retrieval ≠ Generation

High recall (LlamaIndex) does not guarantee better F1 scores, but it significantly improves **semantic alignment** (BARTScore) and answer fidelity.

⚠ **Quality vs. Token Match**



Tool Orchestration Critical

Tool efficiency is a first-class metric. Superior tool selection correlates with better semantic quality, proving that **agent design** directly impacts output reliability.

✖ **Minimize Waste Calls**



Multi-Metric Evaluation is Essential

Single metrics mislead. A robust evaluation must triangulate **F1/BLEU** (text match), **Recall@k** (retrieval), **Latency** (UX), and **Tool Efficiency** (system health) to capture the full picture of RAG performance.

⬇ **Holistic Benchmarking**



Deep Instrumentation Pays Off

Granular tracing (e.g., WandB) is non-negotiable for RAG. It reveals hidden behaviors like **latency tails** at depth k=100 and precise **agent decision paths** that aggregate metrics obscure.

👁 **Visibility = Reliability**

7. Future Directions

Roadmap for Legal RAG Research

 Next Steps

Next Steps

Multi-hop Reasoning Evaluation

Extend benchmark to queries requiring synthesis across multiple contract sections (cross-document reasoning), addressing current single-hop limitations.

Hybrid Architectures

Combine **LlamaIndex**'s superior retrieval backend with **LangChain**'s low-latency orchestration engine to optimize the recall-latency frontier.

Advanced Agentic Tools

Move beyond keyword heuristics to LLM-based tool selection. Implement complex tools for clause comparison and risk scoring.

Domain Extension

Validate findings on high-stakes domains like **Medical (EHR)** and **Finance (10-K)**, varying chunking strategies and embedding models.

Links and references

Slides beautified with Genspark

Github Link: [Link](#)

Wanb Link: [Link](#)

Wand Report Link: [Link](#)

Report Link: [Link](#)

Thank you